

## Basic Checkpoint System-call Design

Our basic design first requires a quick, small change in the kernel task structure. We added a counter variable, `checkpoints`, to `task_struct` which is initialized to 0. This value is incremented during every checkpoint (i.e. every call to the `cp_range()` system call), and thus can be interpreted as the number of checkpoints run on a given process. In addition, we define a simple function, `page_dump(addr)`, which dumps a given page to a file specified by the PID and number of checkpoints taken already. It does so by looping from `addr` to `addr+PAGE_SIZE` (assuming `addr` is the beginning of a page of memory), and writing the value of each memory address in between to a file located at `/tmp/<pid>-cp<n>`, where `pid = current->pid` and `n = current->checkpoints`. This I/O is supported by the `sys_open()`, `sys_write()`, and `sys_close()` functions.

The actual implementation of the system call is relatively straightforward. To begin with, the `vm_area_structs` containing `*start` and `*end` are found by using `vma_find()`. These values are clamped from 0GB-3GB to ensure that the checkpoint will only include user code. From here, we loop through the `vm_next` pointer across all of the found VMAs. For each VMA region, we loop from `vm_start` to `vm_end` by increments of `PAGE_SIZE` if and only if the VMA is not `VM_IO` or `VM_LOCKED` and has `VM_READ` enabled. For each of these addresses, `page_dump()` is called if and only if the location conditions and page status flags are satisfied. The location conditions are an AND expression that determines if the given page start and end addresses contain the `*start` and `*end` addresses. The page status check is implemented in the `page_valid()` function, which ensures that the page is not locked, NULL, errored, or `ZERO_PAGE`. This ensures that valid pages with the beginning or tail end of the memory passed in are dumped as well as valid pages that contain all of the data within, covering all of our needed cases. Again, if this condition is true, the `page_dump()` function then appends the data stored in that page along with the page start address to a file in `/tmp`. Upon calling the system call `cp_range()` again, `current->checkpoints` will be incremented so output will go to the next file.

## Incremental Checkpoint System-call Design

Our incremental design is built on top of our implementation of the basic full checkpointing system. Thus, this includes changes made to `task_struct`, the `cp_range(void *start, void *end)` function, and the `page_dump(addr)` function. It will be stored in the same file and implemented in the `inc_cp_range(void *start, void *end)` function.

The incremental checkpoint process may be split into two cases: the initial checkpoint and the subsequent checkpoints made afterwards. In the case of the initial checkpoint for a

process, `current->checkpoints` is read. If the value is 0, a full checkpoint is needed and is done by simply calling the `cp_range(void *start, void *end)` function from the original implementation in Part 1 of the lab and then immediately returning. This will write a file to `/tmp/<pid>-cp<n>` as before. This is due to the fact that the initial checkpoint must always be a full checkpoint.

For each subsequent checkpoint, additional logic is needed to determine which pages to save. To implement the incremental behavior, we check the dirty bit for each page before deciding whether or not to save it. This bit shows whether or not the page has been modified since being paged in. The file `include/linux/page-flags.h` contains information regarding the flags variable within each page structure. `PG_dirty` is defined as the bit position that holds the dirty bit and the macro `pte_dirty(dirty)` tests it by using `test_bit(bit, value)`. Thus, our incremental checkpoint system call can determine if a page has been modified by using this macro. To implement the incremental case, the system call must loop through each VMA from the `*start` to `*end` address, just as in the original case. For each `vm_area_struct` found, the code must also loop through all addresses from `vma_start` to `vma_end` by increments of `PAGE_SIZE`, storing the page start in an unsigned long address variable. This is equivalent to looping over each page in the VMA. From here, the `follow_page()` function from `mm/memory.c` may be called, yielding a pointer to a page structure given the address. Finally, with a pointer to the page, `pte_dirty(page)` may be called. If the result of this macro is true (non-zero valued), the page has been modified since last page-in and must be saved. From here, the `page_dump(addr)` function (taken from our original implementation) may be called. This opens the file `/tmp/<pid>-cp<n>` and writes out the contents of all memory addresses from `addr` to `addr+PAGE_SIZE`. Thus, only the modified pages will be saved to the output file.

## Special Cases

Following were the checks we added for the pages and the VMAs:

`PG_LOCKED`, `PG_ERROR`, `PG_RECLAIM`, `PG_MAPPEDTODISK`, `PG_SWAPCACHE`, `PG_COMPOUND`, `PG_NOSAVE`, `PG_WRITEBACK`, `PG_PRIVATE`, `PG_CHECKED`, `PG_SLAB`, `PG_ACTIVE`, `PG_LRU`, `PG_UPTODATE`, `PG_REFERENCED`, `PG_HIGHMEM` and `PG_RESERVED`.

Flags checked for VMA -- `VM_LOCKED`, `VM_IO`, `VM_READ`, `VM_DONTCOPY`, `VM_RESERVED`.

As an optimization step, we also check if a pg is a “zero\_page” before taking a backup of it.

## Bugs/Limitations

While taking the full system dump (from address 0 to 3GB), we encountered this error - “Unable to handle kernel paging request at virtual address 003a6000”. We tried debugging it by

printing the VMA values, adding all possible checks for the pages and the VMA areas, but we still could not figure it out. The program takes the backup up until the address 0x003a6000, but stops when it encounters this address.